

# Machine Intelligence:: Deep Learning

## Week 5





*Beate Sick, Elvis Murina, Oliver Dürr*

Institut für Datenanalyse und Prozessdesign  
Zürcher Hochschule für Angewandte Wissenschaften

Part I: Probabilistic models with flexible CPDs

Winterthur, 24. March. 2020

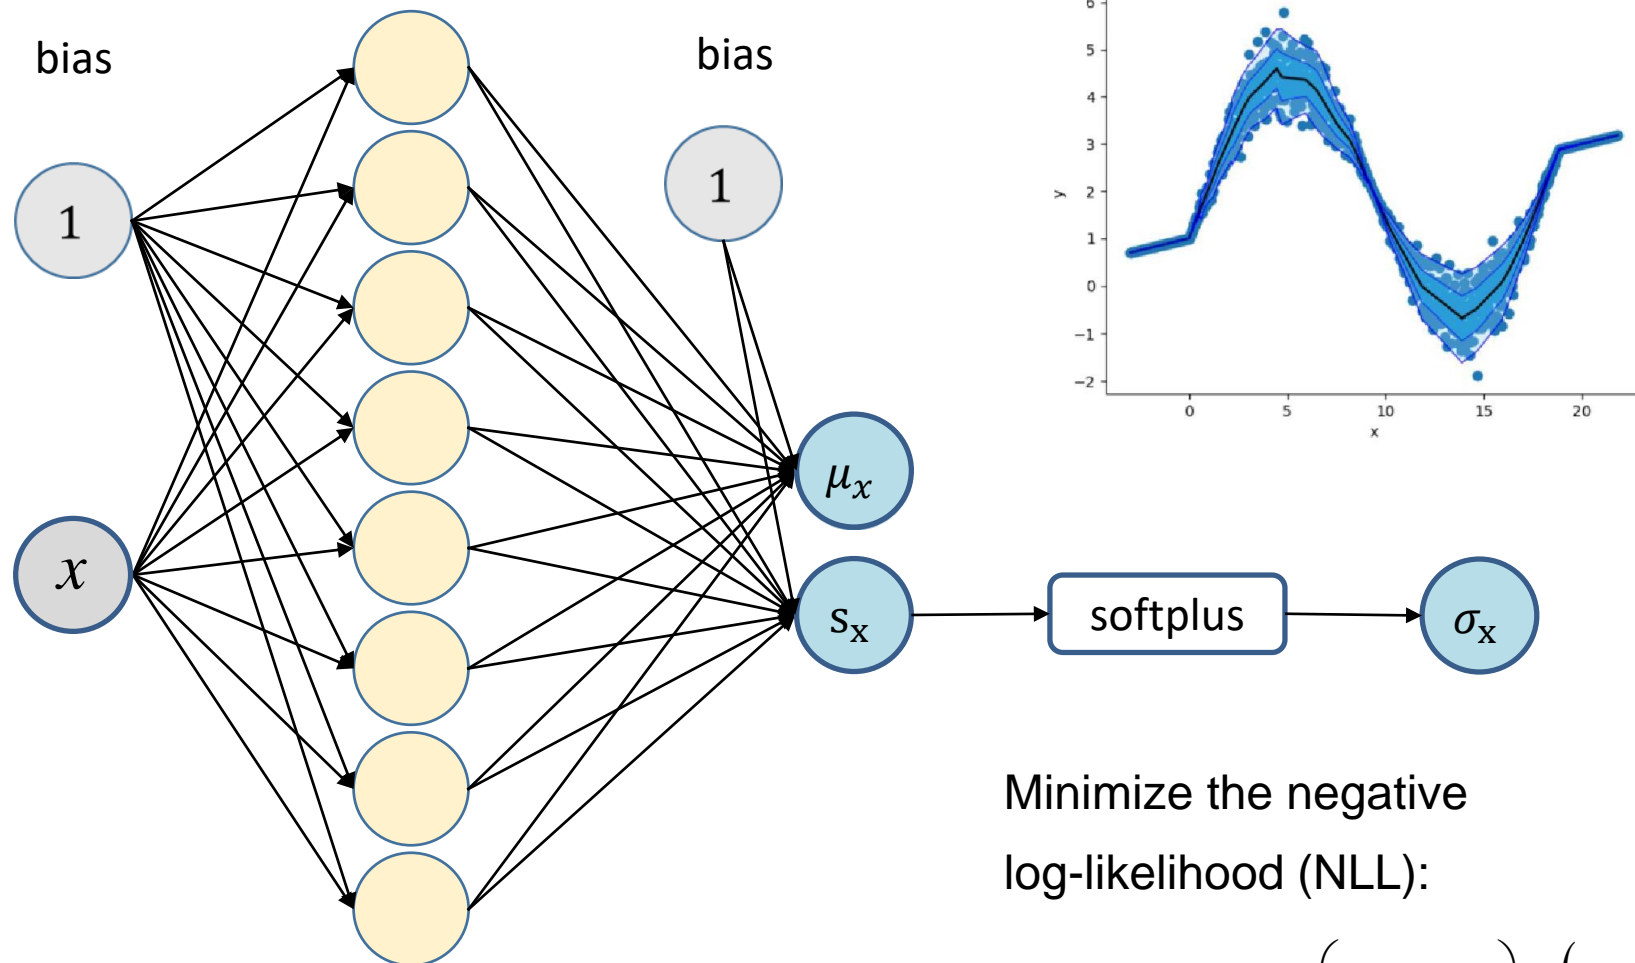
We have a flexible tool where the choice of the architecture and the choice of the outcome distribution is independent



# Modeling continuous data

## continued




# Modeling a flexible variance



Minimize the negative  
log-likelihood (NLL):

$$\hat{\mathbf{w}}_{\text{ML}} = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_{i=1}^n -\log \left( \frac{1}{\sqrt{2\pi\sigma_x^2}} \right) + \frac{(y_i - \mu_{x_i})^2}{2\sigma_{x_i}^2}$$

# Modeling a constant variance



Minimize the negative log-likelihood (NLL):

$$\hat{\mathbf{w}}_{\text{ML}} = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_{i=1}^n -\log \left( \frac{1}{\sqrt{2\pi\sigma^2}} \right) + \frac{(y_i - \mu_{x_i})^2}{2\sigma^2}$$

# Modeling count data continued

# The camper example - tabular data modeled via fcNN

N=250 groups visiting a national park

**Y=count: number of fishes caught**

X1=persons: number of persons in group

X2=child: number of children in the group

X3=bait: indicates if bait was used

X4=camper: indicates if camper is brought





Data: <https://stats.idre.ucla.edu/r/dae/zip>

# Model 1: linear regression, predicted CPDs for test observations

Predict CPD for outcome in test data:


Group 31 used livebait, had a camper and were 4 persons with one child.  $Y=5$  fish.

Group 33 used livebait, didn't have a camper and were 4 persons with two children.  $Y=0$  fish.





# Model 1: linear regression, visualize the CPDs by quantiles



The mean of the CPD is depicted by the solid lines.

The dashed lines represent the 0.025 and 0.975 quantiles, yielding the borders of a 95% prediction interval.



Note that different combinations of predictor values can yield the same parameters of the CPD.

# Model 2: Poisson regression, predicted CPDs for test observations

Predict CPD for outcome in test data:

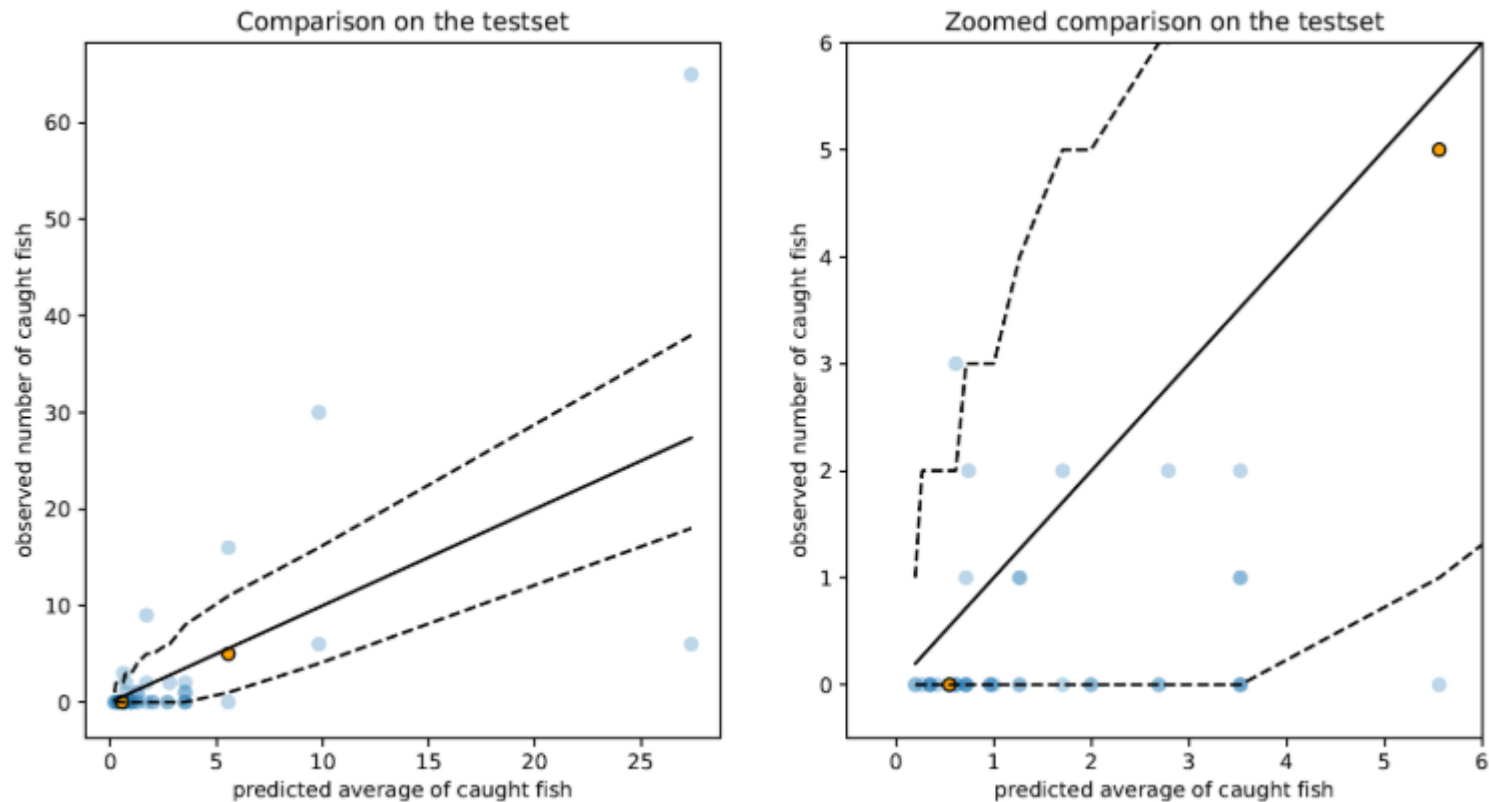
Group 31 used livebait, had a camper and were 4 persons with one child.  $Y=5$  fish.

Group 33 used livebait, didn't have a camper and were 4 persons with two children.  $Y=0$  fish.



What is the likelihood of the observed outcome in test obs 31 and 33?

## Model 2: Poisson regression, visualize the CPDs by quantiles



The mean of the CPD is depicted by the solid lines.

The dashed lines represent the 0.025 and 0.975 quantiles, yielding the borders of a 95% prediction interval.

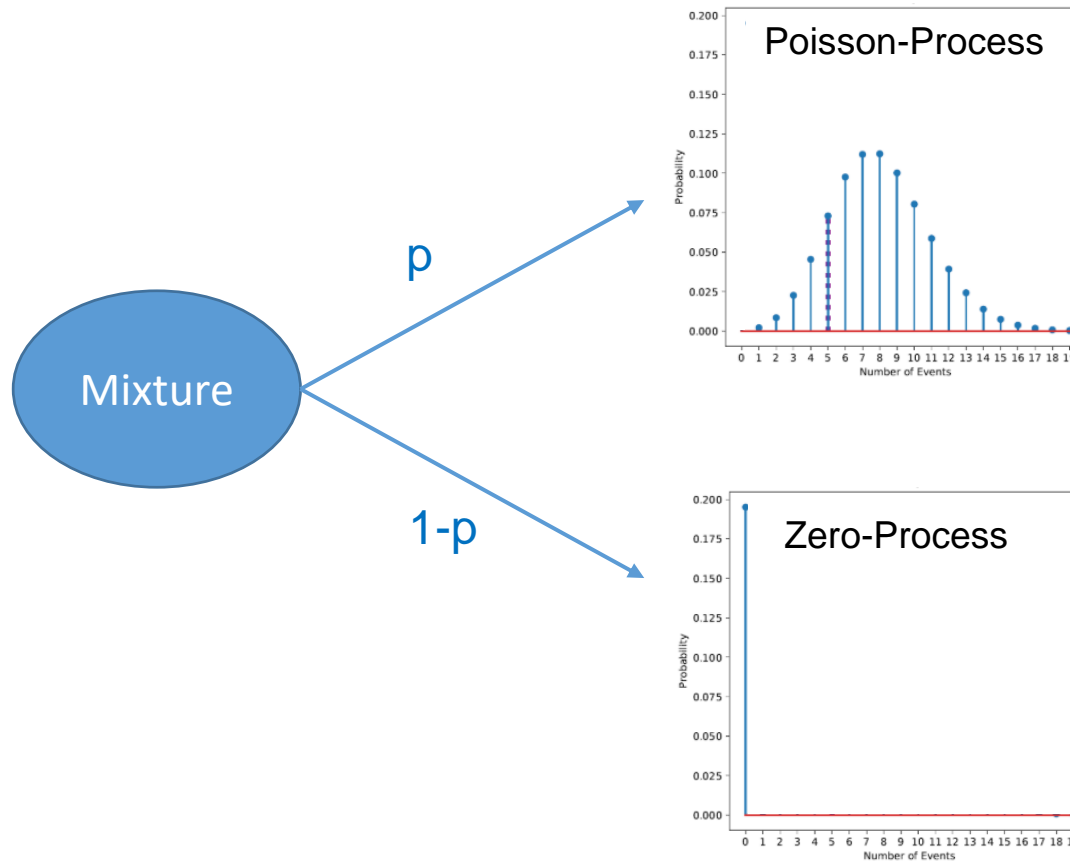
Note that different combinations of predictor values can yield the same parameters of the CPD.

**Modeling count data:**

**M3: ZIP regression**


# Zero-Inflated Poisson (ZIP) as Mixture Process

How many fish a group catches does not only depend on luck ;-)




# Zero-Inflated Poisson (ZIP) can be seen as Mixture Distribution

Poisson-Process




$+(1-p) \cdot$

Zero-Process



$=$

Zero-inflated Poisson



# Custom distribution for a ZIP distribution

```
def zero_inf(out):  
    rate = tf.squeeze(tf.math.exp(out[:,0:1])) # First NN output controls lambda. Exp guarantee value >0  
    s = tf.math.sigmoid(out[:,1:2]) # Second NN output controls p; sigmoid guarantees value in [0,1]  
    probs = tf.concat([1-s, s], axis=1) # The two probabilities for 0's or Poissonian distribution  
    return tfd.Mixture(  
        cat=tfd.Categorical(probs=probs), # tfd.Categorical allows creating a mixture of two components  
        components=[  
            tfd.Deterministic(loc=tf.zeros_like(rate)), # Zero as a deterministic value  
            tfd.Poisson(rate=rate), # Value drawn from a Poissonian distribution  
        ]  
    )
```

## Model 3: Zero-Inflated Poisson regression via NNs in keras

```
## Definition of the custom parameterized distribution  
inputs = tf.keras.layers.Input(shape=(X_train.shape[1],))  
out = Dense(2)(inputs) #A  
p_y_zi = tfp.layers.DistributionLambda(zero_inf)(out)  
model_zi = Model(inputs=inputs, outputs=p_y_zi)
```





# Model 3: ZIP regression, get test NLL from Gaussian CPD

Predict CPD for outcome in test data:


Group 31 used livebait, had a camper and were 4 persons with one child. Y=5 fish.

Group 33 used livebait, didn't have a camper and were 4 persons with two children. Y=0 fish.



What is the likelihood of the observed outcome in test obs 31 and 33?

## Model 3: ZIP regression, visualize the CPDs by quantiles




The mean of the CPD is depicted by the solid lines.

The dashed lines represent the 0.025 and 0.975 quantiles, yielding the borders of a 95% prediction interval.


Note that different combinations of predictor values can yield the same parameters of the CPD.

# Validation NLL allows to rank different probabilistic models



# Probabilistic models with complex input data


# The UTK face data - face image data with known age



Data: <https://stats.idre.ucla.edu/r/dae/zip>

UTKFace data set containing N= 23'708 images of cropped faces of humans with known age ranging from 1 to 116 years.

# Modeling a flexible variance



Minimize the negative log-likelihood (NLL):

$$\hat{\mathbf{w}}_{\text{ML}} = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_{i=1}^n -\log \left( \frac{1}{\sqrt{2\pi\sigma^2}} \right) + \frac{(y_i - \mu_{x_i})^2}{2\sigma^2}$$

# CNNs for modeling Gaussian CPDs

```
def NLL(y, distr):
    return -distr.log_prob(y)

def my_dist(params):
    return tfd.Normal(loc=params[:,0:1], scale=1e-3 + tf.math.softplus(0.05 * params[:,1:2]))# both parameters are learnable

inputs = Input(shape=(80,80,3))
x = Convolution2D(16,kernel_size,padding='same',activation="relu")(inputs)
x = Convolution2D(16,kernel_size,padding='same',activation="relu")(x)
x = MaxPooling2D(pool_size=pool_size)(x)

x = Convolution2D(32,kernel_size,padding='same',activation="relu")(x)
x = Convolution2D(32,kernel_size,padding='same',activation="relu")(x)
x = MaxPooling2D(pool_size=pool_size)(x)



x = Convolution2D(32,kernel_size,padding='same',activation="relu")(x)
x = Convolution2D(32,kernel_size,padding='same',activation="relu")(x)
x = MaxPooling2D(pool_size=pool_size)(x)

x = Flatten()(x)
x = Dense(500,activation="relu")(x)
x = Dropout(0.3)(x)
x = Dense(50,activation="relu")(x)
x = Dropout(0.3)(x)
x = Dense(2)(x)
dist = tf.layers.DistributionLambda(my_dist)(x)

model_flex = Model(inputs=inputs, outputs=dist)
model_flex.compile(tf.keras.optimizers.Adam(), loss=NLL)
```

We model both parameters ( $\mu_x, \sigma_x$ ) of a Gaussian CPD  $N(\mu_x, \sigma_x)$   
→ More flexible than in classical regression where  $\sigma = \text{constant}$

# Modeling a constant variance



Minimize the negative log-likelihood (NLL):

$$\hat{\mathbf{w}}_{\text{ML}} = \underset{w}{\operatorname{argmin}} \sum_{i=1}^n -\log \left( \frac{1}{\sqrt{2\pi\sigma^2}} \right) + \frac{(y_i - \mu_{x_i})^2}{2\sigma^2}$$



# CNNs for modeling Gaussian CPDs

```
def NLL(y, distr):
    return -distr.log_prob(y)

def my_dist(params):
    return tfd.Normal(loc=params[:,0:1], scale=1e-3 + tf.math.softplus(0.05 * params[:,1:2]))# both parameters are learnable

input1 = Input(shape=(80,80,3))
input2 = Input(shape=(1,))
x = Convolution2D(16,kernel_size,padding='same',activation="relu")(input1)
x = Convolution2D(16,kernel_size,padding='same',activation="relu")(x)
x = MaxPooling2D(pool_size=pool_size)(x)

x = Convolution2D(32,kernel_size,padding='same',activation="relu")(x)
x = Convolution2D(32,kernel_size,padding='same',activation="relu")(x)
x = MaxPooling2D(pool_size=pool_size)(x)


x = Convolution2D(32,kernel_size,padding='same',activation="relu")(x)
x = Convolution2D(32,kernel_size,padding='same',activation="relu")(x)
x = MaxPooling2D(pool_size=pool_size)(x)

x = Flatten()(x)
x = Dense(500,activation="relu")(x)
x = Dropout(0.3)(x)
x = Dense(50,activation="relu")(x)
x = Dropout(0.3)(x)
out1 = Dense(1)(x)
out2 = Dense(1)(input2)
params = Concatenate()([out1,out2])
dist = tfp.layers.DistributionLambda(my_dist)(params) #

model_const_sd = Model(inputs=[input1,input2], outputs=dist) ## use a trick with two inputs, input2 is just ones
model_const_sd.compile(tf.keras.optimizers.Adam(), loss=NLL)
```


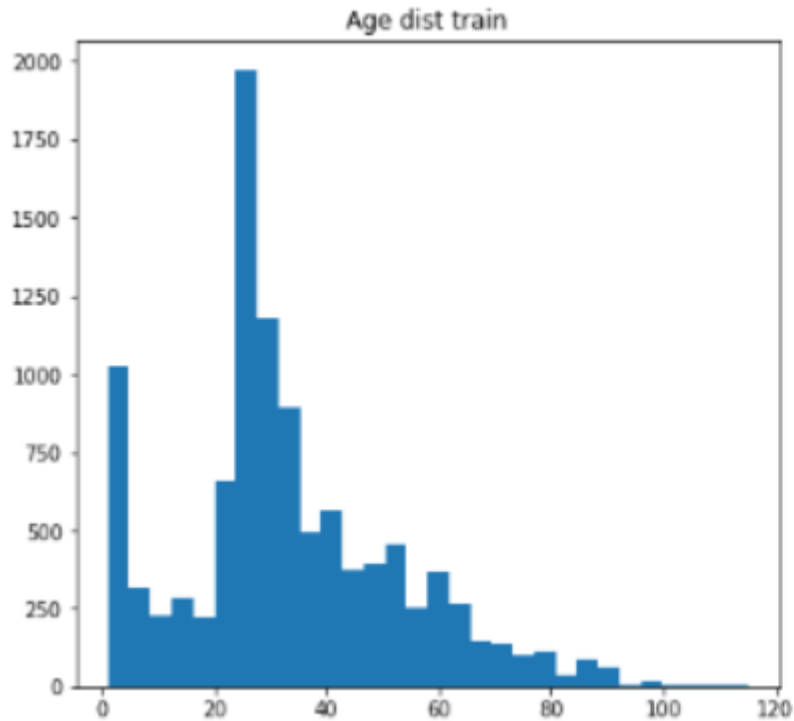
We model both parameters ( $\mu, \sigma$ ) of a Gaussian CPD  $N(\mu_x, \sigma) \rightarrow$  But assume a constant variance

# Excercise



Check out the models for age prediction with flexible and constant variance and try to understand the code and to answer the questions.


# Age distribution




We have a lot of small children in the data set, for whom the age estimation is probably not so difficult.

# Resulting age CPDs

## Constant variance




## flexible variance



In case of a flexible variance, a broad predicted Gaussian CPD does indicate high uncertainty about the age.

# The MNIST data set - images of handwritten digits



The MNIST data set containing  $N = 60'000$  images of handwritten digits of the ten classes 0, 1, 2, 3, 4, 5, 6, 7, 8, 9,

# CNNs for modeling Gaussian CPDs

```
# here we define hyperparameter of the CNN
```

```
batch_size = 128
```

```
nb_classes = 10
```

```
img_rows, img_cols = 28, 28
```

```
kernel_size = (3, 3)
```

```
input_shape = (img_rows, img_cols, 1)
```

```
pool_size = (2, 2)
```

```
# define CNN with 2 convolution blocks and 2 fully connected layers
```

```
model = Sequential()
```

```
model.add(Convolution2D(8, kernel_size, padding='same', input_shape=input_shape))
```

```
model.add(Activation('relu'))
```

```
model.add(Convolution2D(8, kernel_size, padding='same'))
```

```
model.add(Activation('relu'))
```

```
model.add(MaxPooling2D(pool_size=pool_size))
```

```
model.add(Convolution2D(16, kernel_size, padding='same'))
```

```
model.add(Activation('relu'))
```

```
model.add(Convolution2D(16, kernel_size, padding='same'))
```

```
model.add(Activation('relu'))
```

```
model.add(MaxPooling2D(pool_size=pool_size))
```

```
model.add(Flatten())
```

```
model.add(Dense(40))
```

```
model.add(Activation('relu'))
```

```
model.add(Dense(nb_classes))
```

```
model.add(Activation('softmax'))
```

```
# compile model and initialize weights
```


```
model.compile(loss='categorical_crossentropy',
```

```
              optimizer='adam',
```

```
              metrics=['accuracy'])
```

We model 10 parameters ( $p_0, p_1, \dots, p_9$ ) of a Multinomial CPD (it could also be done with 9 parameters)  
→ Most flexible CPD for an outcome with 10 possible values

We predict for each image a multinomial CPD



The multinomial distribution is especially flexible because it has as many parameters as possible values (or actually one parameter less, because probabilities need to sum up to one).